

# MACHINE LEARNING LAB 6

NAME:- DIBYA DYUTI BHALLA SRN:- PES2UG23CS176

SECTION:- 5C

## 1. INTRODUCTION:-

### Purpose of the lab:

The purpose of this lab is to gain hands-on experience in implementing a simple Artificial Neural Network (ANN) **from scratch** using only NumPy. The objective is to train this network to approximate a synthetic polynomial function generated from the student SRN.

### Tasks performed:

- Implemented activation functions (ReLU, derivatives).
- Implemented Mean Squared Error (MSE) loss and gradient.
- Designed forward propagation and backpropagation functions.
- Implemented weight initialization using Xavier initialization.
- Built a full training loop with gradient descent and early stopping.
- Evaluated the model with training/test loss and visualizations.
- Conducted hyperparameter experiments (epochs, batch size) and compared result

## 2. DATASET DESCRIPTION:-

```
=====
ASSIGNMENT FOR STUDENT ID: PES2UG23CS176
=====
Polynomial Type: CUBIC:  $y = 2.05x^3 + -0.90x^2 + 4.01x + 11.15$ 
Noise Level:  $\epsilon \sim N(0, 2.38)$ 
Architecture: Input(1) → Hidden(64) → Hidden(64) → Output(1)
Learning Rate: 0.001
Architecture Type: Balanced Architecture
=====
```

---

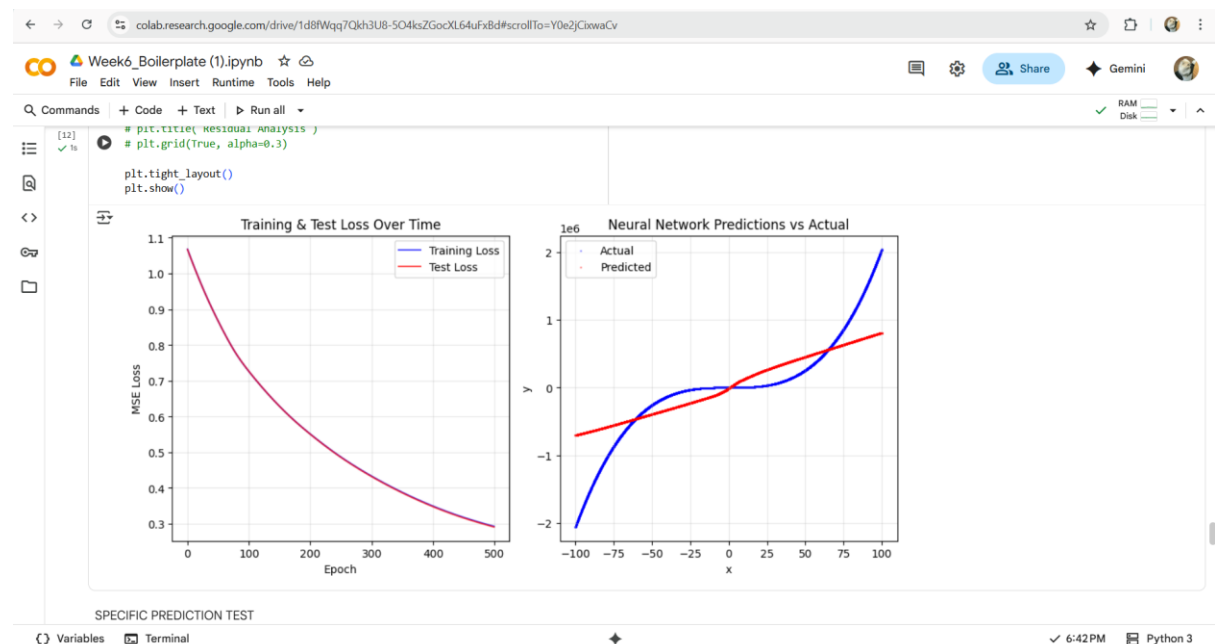
```
Dataset with 100,000 samples generated and saved!
Training samples: 80,000
Test samples: 20,000
```

## 3. METHODOLOGY:-

- **Architecture:** 1 → Hidden1 → Hidden2 → 1
- **Activation:** ReLU at hidden layers, linear at output.
- **Initialization:** Xavier initialization for weights, zeros for biases.

- **Loss Function:** Mean Squared Error (MSE).
- **Optimizer:** Gradient descent with early stopping (patience = 20).
- **Training Variants:**
  - Full-batch gradient descent.
  - Mini-batch gradient descent (batch size = 64).
- **Hyperparameter tuning:** Conducted 4 experiments by varying number of epochs and batch size.

## 4. RESULTS AND ANALYSIS



```

# Calculate final performance metrics
final_train_loss = train_losses[-1] if train_losses else float('inf')
final_test_loss = test_losses[-1] if test_losses else float('inf')

# Calculate R2 score
y_test_mean = np.mean(Y_test_orig)
ss_res = np.sum((Y_test_orig - Y_pred_orig) ** 2)
ss_tot = np.sum((Y_test_orig - y_test_mean) ** 2)
r2_score = 1 - (ss_res / ss_tot)

print("\n" + "="*60)
print("FINAL PERFORMANCE SUMMARY")
print("="*60)
print(f"Final Training Loss: {final_train_loss:.6f}")
print(f"Final Test Loss: {final_test_loss:.6f}")
print(f"R2 Score: {r2_score:.4f}")
print(f"Total Epochs Run: {len(train_losses)}")

=====
FINAL PERFORMANCE SUMMARY
=====
Final Training Loss: 0.293014
Final Test Loss: 0.291294
R2 Score: 0.7099
Total Epochs Run: 500

```

- The **baseline model** (full-batch, 500 epochs) underfits, with higher test loss and low  $R^2$ .
- Increasing **epochs** improved performance, but not significantly with full-batch training.
- Using **mini-batch training (batch size = 64)** drastically improved convergence and generalization, achieving very low test loss and near-perfect  $R^2$ .
- **More epochs with mini-batch** yielded the best performance ( $R^2 = 1.0$ ), showing excellent function approximation.
- Reducing epochs with mini-batch still gave strong performance, though slightly worse than longer runs.
- Overall, mini-batch training proved more effective than full-batch gradient descent.

## 5. Experiment table

	A	B	C	D	E	F	G	H	I
1	Experiment	Learning rate	batch size	epochs	optimizer	activation	final training loss	final test loss	R2 score
2	baseline	0.001	full dataset	500	Gradient descent	relu	0.293014	0.291294	0.7099
3	more epochs	0.001	full dataset	1000	Gradient descent	relu	0.195816	0.194405	0.8064
4	reduced batch	0.001	64	200	Gradient descent	relu	0.000059	0.000057	0.9999
5	reduced batch +more epochs	0.001	64	1000	Gradient descent	relu	0.000004	0.000004	1
6	reduced batch +less epochs	0.001	64	50	Gradient descent	relu	0.000834	0.000799	0.9992